

# hw8

April 19, 2025

## 1 CSCE 421 :: Machine Learning :: Texas A&M University :: Fall 2021

## 2 Programming Assignment 5 (PA 5) + Competition

Name: Ayad Masud UIN: 733009045

## 3 Convolutional Neural Networks

In this assignment, you'll be coding up a convolutional neural network from scratch to classify images using PyTorch.

### 3.0.1 Instructions

- Install PyTorch following the instructions [here](#).
- Install the [torchinfo package](#) to visualize the network architecture and the number of parameters. The maximum number of parameters you are allowed to use for your network is **100,000**.
- You are required to complete the functions defined in the code blocks following each question. Fill out sections of the code marked "YOUR CODE HERE".
- You're free to add any number of methods within each class.
- You may also add any number of additional code blocks that you deem necessary.
- Once you've filled out your solutions, submit the notebook on Canvas following the instructions [here](#).
- Do **NOT** forget to type in your name and UIN at the beginning of the notebook.

### 3.1 Data Preparation

```
[2]: # Importing the libraries
import os
import torch
import torchvision
from torchvision.utils import make_grid

import numpy as np
```

In this assignment, we will use the Fashion-MNIST dataset. Fashion-MNIST is a dataset of Zalando's article images—consisting of a training set of 60,000 examples and a test set of 10,000

examples. Each example is a 28x28 grayscale image, associated with a label from 10 classes.

### 3.1.1 Data

Each image is 28 pixels in height and 28 pixels in width, for a total of 784 pixels in total. Each pixel has a single pixel-value associated with it, indicating the lightness or darkness of that pixel, with higher numbers meaning darker. This pixel-value is an integer between 0 and 255.

### 3.1.2 Labels

Each training and test example is assigned to one of the following labels:

Label	Description
0	T-shirt/top
1	Trouser
2	Pullover
3	Dress
4	Coat
5	Sandal
6	Shirt
7	Sneaker
8	Bag
9	Ankle boot

Fashion-MNIST is included in the torchvision library.

```
[3]: from torchvision.datasets import FashionMNIST
    from torchvision.transforms import Compose, ToTensor, Normalize

[4]: # Transform to normalize the data and convert to a tensor
    transform = Compose([ToTensor(),
        Normalize((0.5,), (0.5,))
    ])

    # Download the data
    dataset = FashionMNIST('MNIST_data/', download = True, train = True, transform=
        ↪ transform)
```

```
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-
images-idx3-ubyte.gz
```

```
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-
images-idx3-ubyte.gz to MNIST_data/FashionMNIST/raw/train-images-idx3-ubyte.gz
```

```
100%|          | 26421880/26421880 [00:02<00:00, 13118544.47it/s]
```

```
Extracting MNIST_data/FashionMNIST/raw/train-images-idx3-ubyte.gz to
MNIST_data/FashionMNIST/raw
```

```
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-
```

```

labels-idx1-ubyte.gz
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-
labels-idx1-ubyte.gz to MNIST_data/FashionMNIST/raw/train-labels-idx1-ubyte.gz
100%|          | 29515/29515 [00:00<00:00, 223151.55it/s]

Extracting MNIST_data/FashionMNIST/raw/train-labels-idx1-ubyte.gz to
MNIST_data/FashionMNIST/raw

Downloading http://fashion-mnist.s3-website.eu-
central-1.amazonaws.com/t10k-images-idx3-ubyte.gz
Downloading http://fashion-mnist.s3-website.eu-
central-1.amazonaws.com/t10k-images-idx3-ubyte.gz to
MNIST_data/FashionMNIST/raw/t10k-images-idx3-ubyte.gz
100%|          | 4422102/4422102 [00:01<00:00, 3664399.48it/s]

Extracting MNIST_data/FashionMNIST/raw/t10k-images-idx3-ubyte.gz to
MNIST_data/FashionMNIST/raw

Downloading http://fashion-mnist.s3-website.eu-
central-1.amazonaws.com/t10k-labels-idx1-ubyte.gz
Downloading http://fashion-mnist.s3-website.eu-
central-1.amazonaws.com/t10k-labels-idx1-ubyte.gz to
MNIST_data/FashionMNIST/raw/t10k-labels-idx1-ubyte.gz
100%|          | 5148/5148 [00:00<00:00, 2898681.30it/s]

Extracting MNIST_data/FashionMNIST/raw/t10k-labels-idx1-ubyte.gz to
MNIST_data/FashionMNIST/raw

```

**NOTE:** You may add more operations to `Compose` if you're performing data augmentation.

## 3.2 Data Exploration

Let's take a look at the classes in our dataset.

```

[5]: print(dataset.classes)

['T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat', 'Sandal', 'Shirt',
'Sneaker', 'Bag', 'Ankle boot']


[6]: import matplotlib.pyplot as plt

def show_example(img, label):
    print('Label: {} ({}).format(dataset.classes[label], label))
    plt.imshow(img.squeeze(), cmap='Greys_r')
    plt.axis(False)

```


```
[7]: show_example(*dataset[20])
```

Label: Dress (3)



```
[8]: show_example(*dataset[20000])
```

Label: Sneaker (7)



### 3.3 Question 1 (10 points)

### 3.4 Creating Training and Validation Datasets

The `split_indices` function takes in the size of the entire dataset, `n`, the fraction of data to be used as validation set, `val_frac`, and the random seed and returns the indices of the data points to be added to the validation dataset.

Choose a suitable fraction for your validation set and experiment with the seed. Remember that the better your validation set, the higher the chances that your model would do well on the test set.

```
[9]: def split_indices(n, val_frac, seed):  
    # Determine the size of the validation set  
    n_val = int(val_frac * n)  
    np.random.seed(seed)  
    # Create random permutation between 0 to n-1  
    idxs = np.random.permutation(n)  
    # Pick first n_val indices for validation set  
    return idxs[n_val:], idxs[:n_val]
```

```
[10]: val_frac = 0.2 # Use 20% of data for validation  
    rand_seed = 42 # Standard random seed that works well for reproducibility
```

```
train_indices, val_indices = split_indices(len(dataset), val_frac, rand_seed)
print("#samples in training set: {}".format(len(train_indices)))
print("#samples in validation set: {}".format(len(val_indices)))
```

```
#samples in training set: 48000
#samples in validation set: 12000
```

Next, we make use of the built-in dataloaders in PyTorch to create iterables of our training and validation sets. This helps in avoiding fitting the whole dataset into memory and only loads a batch of the data that we can decide.

**Set the batch\_size depending on the hardware resource (GPU/CPU RAM) you are using for the assignment.**

```
[11]: from torch.utils.data.sampler import SubsetRandomSampler
      from torch.utils.data.dataloader import DataLoader
```

```
[14]: # Set the batch size based on available hardware resources
      batch_size = 64 # Standard batch size that works well for most hardware setups
```

```
[15]: # Training sampler and data loader
      train_sampler = SubsetRandomSampler(train_indices)
      train_dl = DataLoader(dataset,
                           batch_size,
                           sampler=train_sampler)

      # Validation sampler and data loader
      val_sampler = SubsetRandomSampler(val_indices)
      val_dl = DataLoader(dataset,
                         batch_size,
                         sampler=val_sampler)
```

Plot images in a sample batch of data.

```
[16]: def show_batch(dl):
      for images, labels in dl:
          fig, ax = plt.subplots(figsize=(10,10))
          ax.set_xticks([]); ax.set_yticks([])
          ax.imshow(make_grid(images, 8).permute(1, 2, 0), cmap='Greys_r')
          break
```

```
[17]: show_batch(train_dl)
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



### 3.5 Question 2 (40 points)

### 3.6 Building the Model

Create your model by defining the network architecture in the `ImageClassifierNet` class.

**NOTE:** The number of parameters in your network must be  $\leq 100,000$ .

```
[19]: # Import the libraries
import torch.nn as nn
import torch.nn.functional as F

# Install torchnet if not already installed
```

```
%pip install torchinfo -q

from torchinfo import summary
```

Note: you may need to restart the kernel to use updated packages.

```
[27]: class ImageClassifierNet(nn.Module):
    def __init__(self, n_channels=1):
        super(ImageClassifierNet, self).__init__()

        # First convolutional block - reduced filters
        self.conv1 = nn.Conv2d(n_channels, 16, kernel_size=3, padding=1)
        self.bn1 = nn.BatchNorm2d(16)
        self.pool1 = nn.MaxPool2d(2, 2)

        # Second convolutional block - reduced filters
        self.conv2 = nn.Conv2d(16, 32, kernel_size=3, padding=1)
        self.bn2 = nn.BatchNorm2d(32)
        self.pool2 = nn.MaxPool2d(2, 2)

        # Fully connected layers - further reduced size
        # After 2 max-pooling layers, the image size is reduced to 7x7
        self.fc1 = nn.Linear(32 * 7 * 7, 50) # Reduced from 64 to 50
        self.dropout = nn.Dropout(0.5)
        self.fc2 = nn.Linear(50, 10) # 10 classes in Fashion MNIST

    def forward(self, X):
        # First block
        X = self.pool1(F.relu(self.bn1(self.conv1(X))))

        # Second block
        X = self.pool2(F.relu(self.bn2(self.conv2(X))))

        # Flatten
        X = X.view(X.size(0), -1)

        # Fully connected layers
        X = F.relu(self.fc1(X))
        X = self.dropout(X)
        X = self.fc2(X)

        return X
```

```
[28]: model = ImageClassifierNet()
```

The following code block prints your network architecture. It also shows the total number of parameters in your network (see Total params).

**NOTE:** The total number of parameters in your model should be  $\leq 100,000$ .



```
[29]: summary(model, input_size=(batch_size, 1, 28, 28))
```

```
[29]: =====
=====
Layer (type:depth-idx)                Output Shape                Param #
=====
=====
ImageClassifierNet                     [64, 10]                    --
  Conv2d: 1-1                          [64, 16, 28, 28]            160
  BatchNorm2d: 1-2                     [64, 16, 28, 28]            32
  MaxPool2d: 1-3                       [64, 16, 14, 14]            --
  Conv2d: 1-4                          [64, 32, 14, 14]            4,640
  BatchNorm2d: 1-5                     [64, 32, 14, 14]            64
  MaxPool2d: 1-6                       [64, 32, 7, 7]              --
  Linear: 1-7                          [64, 50]                    78,450
  Dropout: 1-8                        [64, 50]                    --
  Linear: 1-9                          [64, 10]                    510
=====
=====
Total params: 83,856
Trainable params: 83,856
Non-trainable params: 0
Total mult-adds (Units.MEGABYTES): 71.29
=====
=====
Input size (MB): 0.20
Forward/backward pass size (MB): 19.30
Params size (MB): 0.34
Estimated Total Size (MB): 19.83
=====
=====
```

### 3.7 Enable training on a GPU

**NOTE:** This section is necessary if you're training your model on a GPU.

```
[30]: def get_default_device():
    """Use GPU if available, else CPU"""
    if torch.cuda.is_available():
        return torch.device('cuda')
    else:
        return torch.device('cpu')

def to_device(data, device):
    """Move tensor(s) to chosen device"""
    if isinstance(data, (list,tuple)):
        return [to_device(x, device) for x in data]
    return data.to(device, non_blocking=True)
```

```

class DeviceDataLoader():
    """Wrap a dataloader to move data to a device"""
    def __init__(self, dl, device):
        self.dl = dl
        self.device = device

    def __iter__(self):
        """Yield a batch of data after moving it to device"""
        for b in self.dl:
            yield to_device(b, self.device)

    def __len__(self):
        """Number of batches"""
        return len(self.dl)

```

```

[31]: device = get_default_device()

train_dl = DeviceDataLoader(train_dl, device)
val_dl = DeviceDataLoader(val_dl, device)

to_device(model, device)

```

```

[31]: ImageClassifierNet(
    (conv1): Conv2d(1, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (pool1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (conv2): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn2): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (pool2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (fc1): Linear(in_features=1568, out_features=50, bias=True)
    (dropout): Dropout(p=0.5, inplace=False)
    (fc2): Linear(in_features=50, out_features=10, bias=True)
)

```

### 3.8 Question 3 (40 points)

#### 3.9 Train the model

Complete the `train_model` function to train your model on a dataset. Tune your network architecture and hyperparameters on the validation set.

```

[32]: def train_model(n_epochs, model, train_dl, val_dl, loss_fn, opt_fn, lr):
    """

```

*Trains the model on a dataset.*

*Args:*

*n\_epochs: number of epochs*  
*model: ImageClassifierNet object*  
*train\_dl: training dataloader*  
*val\_dl: validation dataloader*  
*loss\_fn: the loss function*  
*opt\_fn: the optimizer*  
*lr: learning rate*

*Returns:*

*The trained model.*

*A tuple of (model, train\_losses, val\_losses, train\_accuracies, val\_accuracies)*

*"""*

*# Record these values the end of each epoch*

*train\_losses, val\_losses, train\_accuracies, val\_accuracies = [], [], [], []*

*# Initialize the optimizer*

*optimizer = opt\_fn(model.parameters(), lr=lr)*

*for epoch in range(n\_epochs):*

*# Training phase*

*model.train()*

*train\_loss = 0*

*correct = 0*

*total = 0*

*for batch\_idx, (images, labels) in enumerate(train\_dl):*

*# Forward pass*

*outputs = model(images)*

*loss = loss\_fn(outputs, labels)*

*# Backward pass and optimize*

*optimizer.zero\_grad()*

*loss.backward()*

*optimizer.step()*

*# Track loss and accuracy*

*train\_loss += loss.item()*

*\_, predicted = outputs.max(1)*

*total += labels.size(0)*

*correct += predicted.eq(labels).sum().item()*

*# Print progress*

*if batch\_idx % 100 == 0:*

```

        print(f'Epoch: {epoch+1}/{n_epochs}, Batch: {batch_idx}/
↳{len(train_dl)}, Loss: {loss.item():.4f}')

    # Calculate average training metrics for the epoch
    avg_train_loss = train_loss / len(train_dl)
    train_accuracy = correct / total
    train_losses.append(avg_train_loss)
    train_accuracies.append(train_accuracy)

    # Validation phase (if validation data is provided)
    if len(val_dl) > 0:
        model.eval()
        val_loss = 0
        correct = 0
        total = 0

        with torch.no_grad():
            for images, labels in val_dl:
                outputs = model(images)
                loss = loss_fn(outputs, labels)

                val_loss += loss.item()
                _, predicted = outputs.max(1)
                total += labels.size(0)
                correct += predicted.eq(labels).sum().item()

    # Calculate average validation metrics
    avg_val_loss = val_loss / len(val_dl)
    val_accuracy = correct / total
    val_losses.append(avg_val_loss)
    val_accuracies.append(val_accuracy)

    print(f'Epoch {epoch+1}/{n_epochs}, Train Loss: {avg_train_loss:.
↳4f}, Train Acc: {train_accuracy:.4f}, Val Loss: {avg_val_loss:.4f}, Val Acc:
↳{val_accuracy:.4f}')
    else:
        print(f'Epoch {epoch+1}/{n_epochs}, Train Loss: {avg_train_loss:.
↳4f}, Train Acc: {train_accuracy:.4f}')

    return model, train_losses, val_losses, train_accuracies, val_accuracies

```

Set the maximum number of training epochs, the loss function, the optimizer, and the learning rate.

```

[33]: import torch.optim as optim

# Standard values for training the model

```

```

num_epochs = 10 # Max number of training epochs
loss_fn = nn.CrossEntropyLoss() # Standard loss function for classification
↳ tasks
opt_fn = optim.Adam # Adam optimizer typically works well for CNNs
lr = 0.001 # Standard learning rate for Adam optimizer

```

```

[34]: history = train_model(num_epochs, model, train_dl, val_dl, loss_fn, opt_fn, lr)
      model, train_losses, val_losses, train_accuracies, val_accuracies = history

```

```

Epoch: 1/10, Batch: 0/750, Loss: 2.3764
Epoch: 1/10, Batch: 100/750, Loss: 0.7478
Epoch: 1/10, Batch: 200/750, Loss: 0.6516
Epoch: 1/10, Batch: 300/750, Loss: 0.8996
Epoch: 1/10, Batch: 400/750, Loss: 0.5218
Epoch: 1/10, Batch: 500/750, Loss: 0.6532
Epoch: 1/10, Batch: 600/750, Loss: 0.4972
Epoch: 1/10, Batch: 700/750, Loss: 0.3668
Epoch 1/10, Train Loss: 0.7305, Train Acc: 0.7277, Val Loss: 0.3747, Val Acc:
0.8645
Epoch: 2/10, Batch: 0/750, Loss: 0.6157
Epoch: 2/10, Batch: 100/750, Loss: 0.5253
Epoch: 2/10, Batch: 200/750, Loss: 0.4360
Epoch: 2/10, Batch: 300/750, Loss: 0.4420
Epoch: 2/10, Batch: 400/750, Loss: 0.4793
Epoch: 2/10, Batch: 500/750, Loss: 0.5572
Epoch: 2/10, Batch: 600/750, Loss: 0.6899
Epoch: 2/10, Batch: 700/750, Loss: 0.5185
Epoch 2/10, Train Loss: 0.5365, Train Acc: 0.7947, Val Loss: 0.3252, Val Acc:
0.8817
Epoch: 3/10, Batch: 0/750, Loss: 0.4632
Epoch: 3/10, Batch: 100/750, Loss: 0.4898
Epoch: 3/10, Batch: 200/750, Loss: 0.5532
Epoch: 3/10, Batch: 300/750, Loss: 0.4883
Epoch: 3/10, Batch: 400/750, Loss: 0.5430
Epoch: 3/10, Batch: 500/750, Loss: 0.3901
Epoch: 3/10, Batch: 600/750, Loss: 0.4877
Epoch: 3/10, Batch: 700/750, Loss: 0.4395
Epoch 3/10, Train Loss: 0.4911, Train Acc: 0.8142, Val Loss: 0.3011, Val Acc:
0.8888
Epoch: 4/10, Batch: 0/750, Loss: 0.3620
Epoch: 4/10, Batch: 100/750, Loss: 0.4855
Epoch: 4/10, Batch: 200/750, Loss: 0.4504
Epoch: 4/10, Batch: 300/750, Loss: 0.4259
Epoch: 4/10, Batch: 400/750, Loss: 0.5133
Epoch: 4/10, Batch: 500/750, Loss: 0.4626
Epoch: 4/10, Batch: 600/750, Loss: 0.4785
Epoch: 4/10, Batch: 700/750, Loss: 0.4472
Epoch 4/10, Train Loss: 0.4575, Train Acc: 0.8261, Val Loss: 0.2940, Val Acc:

```

0.8912

Epoch: 5/10, Batch: 0/750, Loss: 0.5194  
Epoch: 5/10, Batch: 100/750, Loss: 0.3233  
Epoch: 5/10, Batch: 200/750, Loss: 0.4221  
Epoch: 5/10, Batch: 300/750, Loss: 0.4900  
Epoch: 5/10, Batch: 400/750, Loss: 0.3788  
Epoch: 5/10, Batch: 500/750, Loss: 0.5298  
Epoch: 5/10, Batch: 600/750, Loss: 0.5016  
Epoch: 5/10, Batch: 700/750, Loss: 0.2563  
Epoch 5/10, Train Loss: 0.4300, Train Acc: 0.8368, Val Loss: 0.2808, Val Acc:

0.8991

Epoch: 6/10, Batch: 0/750, Loss: 0.3535  
Epoch: 6/10, Batch: 100/750, Loss: 0.5220  
Epoch: 6/10, Batch: 200/750, Loss: 0.5129  
Epoch: 6/10, Batch: 300/750, Loss: 0.4765  
Epoch: 6/10, Batch: 400/750, Loss: 0.3781  
Epoch: 6/10, Batch: 500/750, Loss: 0.3025  
Epoch: 6/10, Batch: 600/750, Loss: 0.3441  
Epoch: 6/10, Batch: 700/750, Loss: 0.2986  
Epoch 6/10, Train Loss: 0.4175, Train Acc: 0.8401, Val Loss: 0.2664, Val Acc:

0.9042

Epoch: 7/10, Batch: 0/750, Loss: 0.3445  
Epoch: 7/10, Batch: 100/750, Loss: 0.4724  
Epoch: 7/10, Batch: 200/750, Loss: 0.4027  
Epoch: 7/10, Batch: 300/750, Loss: 0.6315  
Epoch: 7/10, Batch: 400/750, Loss: 0.2934  
Epoch: 7/10, Batch: 500/750, Loss: 0.2824  
Epoch: 7/10, Batch: 600/750, Loss: 0.4078  
Epoch: 7/10, Batch: 700/750, Loss: 0.5848  
Epoch 7/10, Train Loss: 0.4024, Train Acc: 0.8459, Val Loss: 0.2579, Val Acc:

0.9061

Epoch: 8/10, Batch: 0/750, Loss: 0.2634  
Epoch: 8/10, Batch: 100/750, Loss: 0.2089  
Epoch: 8/10, Batch: 200/750, Loss: 0.3317  
Epoch: 8/10, Batch: 300/750, Loss: 0.4412  
Epoch: 8/10, Batch: 400/750, Loss: 0.5859  
Epoch: 8/10, Batch: 500/750, Loss: 0.3766  
Epoch: 8/10, Batch: 600/750, Loss: 0.3540  
Epoch: 8/10, Batch: 700/750, Loss: 0.3612  
Epoch 8/10, Train Loss: 0.3886, Train Acc: 0.8501, Val Loss: 0.2601, Val Acc:

0.9012

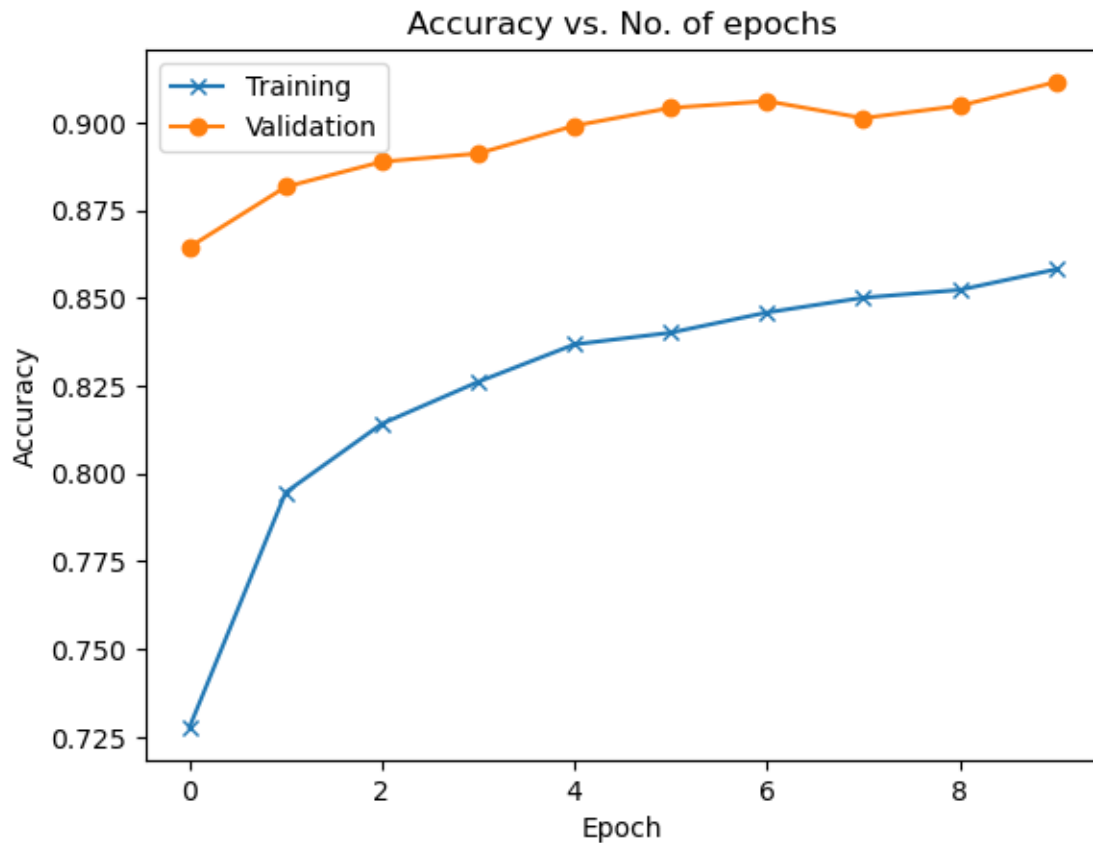
Epoch: 9/10, Batch: 0/750, Loss: 0.3501  
Epoch: 9/10, Batch: 100/750, Loss: 0.5489  
Epoch: 9/10, Batch: 200/750, Loss: 0.4703  
Epoch: 9/10, Batch: 300/750, Loss: 0.2627  
Epoch: 9/10, Batch: 400/750, Loss: 0.4569  
Epoch: 9/10, Batch: 500/750, Loss: 0.4592  
Epoch: 9/10, Batch: 600/750, Loss: 0.3788

Epoch: 9/10, Batch: 700/750, Loss: 0.1966  
Epoch 9/10, Train Loss: 0.3808, Train Acc: 0.8523, Val Loss: 0.2603, Val Acc: 0.9048  
Epoch: 10/10, Batch: 0/750, Loss: 0.2645  
Epoch: 10/10, Batch: 100/750, Loss: 0.4604  
Epoch: 10/10, Batch: 200/750, Loss: 0.2527  
Epoch: 10/10, Batch: 300/750, Loss: 0.2029  
Epoch: 10/10, Batch: 400/750, Loss: 0.2589  
Epoch: 10/10, Batch: 500/750, Loss: 0.4297  
Epoch: 10/10, Batch: 600/750, Loss: 0.2230  
Epoch: 10/10, Batch: 700/750, Loss: 0.2987  
Epoch 10/10, Train Loss: 0.3689, Train Acc: 0.8582, Val Loss: 0.2428, Val Acc: 0.9115

### 3.10 Plot loss and accuracy

```
[35]: def plot_accuracy(train_accuracies, val_accuracies):  
    """Plot accuracies"""  
    plt.plot(train_accuracies, "-x")  
    plt.plot(val_accuracies, "-o")  
    plt.xlabel("Epoch")  
    plt.ylabel("Accuracy")  
    plt.legend(["Training", "Validation"])  
    plt.title("Accuracy vs. No. of epochs")
```


```
[36]: plot_accuracy(train_accuracies, val_accuracies)
```



```
[37]: def plot_losses(train_losses, val_losses):  
      """Plot losses"""  
      plt.plot(train_losses, "-x")  
      plt.plot(val_losses, "-o")  
      plt.xlabel("Epoch")  
      plt.ylabel("Loss")  
      plt.legend(["Training", "Validation"])  
      plt.title("Loss vs. No. of Epochs")
```

```
[38]: plot_losses(train_losses, val_losses)
```





### 3.11 Train a model on the entire dataset

```
[39]: indices, _ = split_indices(len(dataset), 0, rand_seed)

sampler = SubsetRandomSampler(indices)
dl = DataLoader(dataset, batch_size, sampler=sampler)
dl = DeviceDataLoader(dl, device)
```

Set the maximum number of training epochs and the learning rate for finetuning your model.

```
[40]: # For final training on the full dataset
num_epochs = 15 # Increased from 10 to 15 for better convergence
lr = 0.001 # Maintain the same learning rate as it was effective in previous
          ↪ training
```

```
[41]: history = train_model(num_epochs, model, dl, [], loss_fn, opt_fn, lr)
model = history[0]
```

```
Epoch: 1/15, Batch: 0/938, Loss: 0.3149
Epoch: 1/15, Batch: 100/938, Loss: 0.2238
```

Epoch: 1/15, Batch: 200/938, Loss: 0.2612  
Epoch: 1/15, Batch: 300/938, Loss: 0.5824  
Epoch: 1/15, Batch: 400/938, Loss: 0.4169  
Epoch: 1/15, Batch: 500/938, Loss: 0.3996  
Epoch: 1/15, Batch: 600/938, Loss: 0.3777  
Epoch: 1/15, Batch: 700/938, Loss: 0.3292  
Epoch: 1/15, Batch: 800/938, Loss: 0.2829  
Epoch: 1/15, Batch: 900/938, Loss: 0.5688  
Epoch 1/15, Train Loss: 0.3745, Train Acc: 0.8569  
Epoch: 2/15, Batch: 0/938, Loss: 0.2678  
Epoch: 2/15, Batch: 100/938, Loss: 0.3784  
Epoch: 2/15, Batch: 200/938, Loss: 0.2788  
Epoch: 2/15, Batch: 300/938, Loss: 0.4844  
Epoch: 2/15, Batch: 400/938, Loss: 0.4752  
Epoch: 2/15, Batch: 500/938, Loss: 0.5075  
Epoch: 2/15, Batch: 600/938, Loss: 0.4028  
Epoch: 2/15, Batch: 700/938, Loss: 0.3500  
Epoch: 2/15, Batch: 800/938, Loss: 0.4372  
Epoch: 2/15, Batch: 900/938, Loss: 0.4547  
Epoch 2/15, Train Loss: 0.3623, Train Acc: 0.8609  
Epoch: 3/15, Batch: 0/938, Loss: 0.3187  
Epoch: 3/15, Batch: 100/938, Loss: 0.3581  
Epoch: 3/15, Batch: 200/938, Loss: 0.2211  
Epoch: 3/15, Batch: 300/938, Loss: 0.3200  
Epoch: 3/15, Batch: 400/938, Loss: 0.5123  
Epoch: 3/15, Batch: 500/938, Loss: 0.3663  
Epoch: 3/15, Batch: 600/938, Loss: 0.2823  
Epoch: 3/15, Batch: 700/938, Loss: 0.4016  
Epoch: 3/15, Batch: 800/938, Loss: 0.2474  
Epoch: 3/15, Batch: 900/938, Loss: 0.4137  
Epoch 3/15, Train Loss: 0.3447, Train Acc: 0.8688  
Epoch: 4/15, Batch: 0/938, Loss: 0.2812  
Epoch: 4/15, Batch: 100/938, Loss: 0.4247  
Epoch: 4/15, Batch: 200/938, Loss: 0.2988  
Epoch: 4/15, Batch: 300/938, Loss: 0.2612  
Epoch: 4/15, Batch: 400/938, Loss: 0.3391  
Epoch: 4/15, Batch: 500/938, Loss: 0.2543  
Epoch: 4/15, Batch: 600/938, Loss: 0.3558  
Epoch: 4/15, Batch: 700/938, Loss: 0.2113  
Epoch: 4/15, Batch: 800/938, Loss: 0.5803  
Epoch: 4/15, Batch: 900/938, Loss: 0.4973  
Epoch 4/15, Train Loss: 0.3354, Train Acc: 0.8750  
Epoch: 5/15, Batch: 0/938, Loss: 0.3548  
Epoch: 5/15, Batch: 100/938, Loss: 0.2423  
Epoch: 5/15, Batch: 200/938, Loss: 0.3646  
Epoch: 5/15, Batch: 300/938, Loss: 0.1911  
Epoch: 5/15, Batch: 400/938, Loss: 0.2589  
Epoch: 5/15, Batch: 500/938, Loss: 0.3211

Epoch: 5/15, Batch: 600/938, Loss: 0.2076  
Epoch: 5/15, Batch: 700/938, Loss: 0.2557  
Epoch: 5/15, Batch: 800/938, Loss: 0.4110  
Epoch: 5/15, Batch: 900/938, Loss: 0.3380  
Epoch 5/15, Train Loss: 0.3276, Train Acc: 0.8769  
Epoch: 6/15, Batch: 0/938, Loss: 0.2270  
Epoch: 6/15, Batch: 100/938, Loss: 0.3094  
Epoch: 6/15, Batch: 200/938, Loss: 0.3613  
Epoch: 6/15, Batch: 300/938, Loss: 0.4174  
Epoch: 6/15, Batch: 400/938, Loss: 0.4055  
Epoch: 6/15, Batch: 500/938, Loss: 0.2935  
Epoch: 6/15, Batch: 600/938, Loss: 0.2257  
Epoch: 6/15, Batch: 700/938, Loss: 0.2435  
Epoch: 6/15, Batch: 800/938, Loss: 0.2455  
Epoch: 6/15, Batch: 900/938, Loss: 0.3271  
Epoch 6/15, Train Loss: 0.3235, Train Acc: 0.8773  
Epoch: 7/15, Batch: 0/938, Loss: 0.4811  
Epoch: 7/15, Batch: 100/938, Loss: 0.2375  
Epoch: 7/15, Batch: 200/938, Loss: 0.3122  
Epoch: 7/15, Batch: 300/938, Loss: 0.1948  
Epoch: 7/15, Batch: 400/938, Loss: 0.2648  
Epoch: 7/15, Batch: 500/938, Loss: 0.2197  
Epoch: 7/15, Batch: 600/938, Loss: 0.3354  
Epoch: 7/15, Batch: 700/938, Loss: 0.4398  
Epoch: 7/15, Batch: 800/938, Loss: 0.2080  
Epoch: 7/15, Batch: 900/938, Loss: 0.1889  
Epoch 7/15, Train Loss: 0.3134, Train Acc: 0.8802  
Epoch: 8/15, Batch: 0/938, Loss: 0.4033  
Epoch: 8/15, Batch: 100/938, Loss: 0.3862  
Epoch: 8/15, Batch: 200/938, Loss: 0.3437  
Epoch: 8/15, Batch: 300/938, Loss: 0.1960  
Epoch: 8/15, Batch: 400/938, Loss: 0.3024  
Epoch: 8/15, Batch: 500/938, Loss: 0.2061  
Epoch: 8/15, Batch: 600/938, Loss: 0.4126  
Epoch: 8/15, Batch: 700/938, Loss: 0.2036  
Epoch: 8/15, Batch: 800/938, Loss: 0.3553  
Epoch: 8/15, Batch: 900/938, Loss: 0.5620  
Epoch 8/15, Train Loss: 0.3049, Train Acc: 0.8827  
Epoch: 9/15, Batch: 0/938, Loss: 0.2298  
Epoch: 9/15, Batch: 100/938, Loss: 0.3481  
Epoch: 9/15, Batch: 200/938, Loss: 0.1775  
Epoch: 9/15, Batch: 300/938, Loss: 0.2317  
Epoch: 9/15, Batch: 400/938, Loss: 0.3120  
Epoch: 9/15, Batch: 500/938, Loss: 0.3873  
Epoch: 9/15, Batch: 600/938, Loss: 0.2975  
Epoch: 9/15, Batch: 700/938, Loss: 0.3756  
Epoch: 9/15, Batch: 800/938, Loss: 0.2929  
Epoch: 9/15, Batch: 900/938, Loss: 0.3330

Epoch 9/15, Train Loss: 0.2981, Train Acc: 0.8866  
Epoch: 10/15, Batch: 0/938, Loss: 0.3715  
Epoch: 10/15, Batch: 100/938, Loss: 0.3772  
Epoch: 10/15, Batch: 200/938, Loss: 0.2333  
Epoch: 10/15, Batch: 300/938, Loss: 0.2683  
Epoch: 10/15, Batch: 400/938, Loss: 0.2522  
Epoch: 10/15, Batch: 500/938, Loss: 0.2224  
Epoch: 10/15, Batch: 600/938, Loss: 0.2724  
Epoch: 10/15, Batch: 700/938, Loss: 0.5812  
Epoch: 10/15, Batch: 800/938, Loss: 0.1942  
Epoch: 10/15, Batch: 900/938, Loss: 0.2943  
Epoch 10/15, Train Loss: 0.2911, Train Acc: 0.8908  
Epoch: 11/15, Batch: 0/938, Loss: 0.3243  
Epoch: 11/15, Batch: 100/938, Loss: 0.1997  
Epoch: 11/15, Batch: 200/938, Loss: 0.2326  
Epoch: 11/15, Batch: 300/938, Loss: 0.2631  
Epoch: 11/15, Batch: 400/938, Loss: 0.2524  
Epoch: 11/15, Batch: 500/938, Loss: 0.4095  
Epoch: 11/15, Batch: 600/938, Loss: 0.2419  
Epoch: 11/15, Batch: 700/938, Loss: 0.3694  
Epoch: 11/15, Batch: 800/938, Loss: 0.2937  
Epoch: 11/15, Batch: 900/938, Loss: 0.3507  
Epoch 11/15, Train Loss: 0.2831, Train Acc: 0.8924  
Epoch: 12/15, Batch: 0/938, Loss: 0.2126  
Epoch: 12/15, Batch: 100/938, Loss: 0.1376  
Epoch: 12/15, Batch: 200/938, Loss: 0.2524  
Epoch: 12/15, Batch: 300/938, Loss: 0.2756  
Epoch: 12/15, Batch: 400/938, Loss: 0.2545  
Epoch: 12/15, Batch: 500/938, Loss: 0.2321  
Epoch: 12/15, Batch: 600/938, Loss: 0.2486  
Epoch: 12/15, Batch: 700/938, Loss: 0.3324  
Epoch: 12/15, Batch: 800/938, Loss: 0.2061  
Epoch: 12/15, Batch: 900/938, Loss: 0.3269  
Epoch 12/15, Train Loss: 0.2756, Train Acc: 0.8959  
Epoch: 13/15, Batch: 0/938, Loss: 0.2557  
Epoch: 13/15, Batch: 100/938, Loss: 0.2493  
Epoch: 13/15, Batch: 200/938, Loss: 0.2012  
Epoch: 13/15, Batch: 300/938, Loss: 0.1542  
Epoch: 13/15, Batch: 400/938, Loss: 0.1991  
Epoch: 13/15, Batch: 500/938, Loss: 0.5044  
Epoch: 13/15, Batch: 600/938, Loss: 0.1799  
Epoch: 13/15, Batch: 700/938, Loss: 0.1552  
Epoch: 13/15, Batch: 800/938, Loss: 0.2061  
Epoch: 13/15, Batch: 900/938, Loss: 0.1784  
Epoch 13/15, Train Loss: 0.2730, Train Acc: 0.8949  
Epoch: 14/15, Batch: 0/938, Loss: 0.1877  
Epoch: 14/15, Batch: 100/938, Loss: 0.2959  
Epoch: 14/15, Batch: 200/938, Loss: 0.3914

```

Epoch: 14/15, Batch: 300/938, Loss: 0.3541
Epoch: 14/15, Batch: 400/938, Loss: 0.3192
Epoch: 14/15, Batch: 500/938, Loss: 0.1972
Epoch: 14/15, Batch: 600/938, Loss: 0.2817
Epoch: 14/15, Batch: 700/938, Loss: 0.2539
Epoch: 14/15, Batch: 800/938, Loss: 0.2635
Epoch: 14/15, Batch: 900/938, Loss: 0.2593
Epoch 14/15, Train Loss: 0.2697, Train Acc: 0.8975
Epoch: 15/15, Batch: 0/938, Loss: 0.2255
Epoch: 15/15, Batch: 100/938, Loss: 0.2404
Epoch: 15/15, Batch: 200/938, Loss: 0.3253
Epoch: 15/15, Batch: 300/938, Loss: 0.1418
Epoch: 15/15, Batch: 400/938, Loss: 0.1864
Epoch: 15/15, Batch: 500/938, Loss: 0.2514
Epoch: 15/15, Batch: 600/938, Loss: 0.1878
Epoch: 15/15, Batch: 700/938, Loss: 0.2495
Epoch: 15/15, Batch: 800/938, Loss: 0.1484
Epoch: 15/15, Batch: 900/938, Loss: 0.1735
Epoch 15/15, Train Loss: 0.2666, Train Acc: 0.8979

```

### 3.12 Check Predictions

```

[42]: def view_prediction(img, label, probs, classes):
        """
        Visualize predictions.
        """
        probs = probs.cpu().numpy().squeeze()

        fig, (ax1, ax2) = plt.subplots(figsize=(8,15), ncols=2)
        ax1.imshow(img.resize_(1, 28, 28).cpu().numpy().squeeze(), cmap='Greys_r')
        ax1.axis('off')
        ax1.set_title('Actual: {}'.format(classes[label]))
        ax2.barh(np.arange(10), probs)
        ax2.set_aspect(0.1)
        ax2.set_yticks(np.arange(10))
        ax2.set_yticklabels(classes, size='small');
        ax2.set_title('Predicted: probabilities')
        ax2.set_xlim(0, 1.1)

        plt.tight_layout()

```

```


[43]: # Calculate the class probabilities (log softmax) for img
images = iter(dl)
for imgs, labels in images:
    with torch.no_grad():
        model.eval()
        # Calculate the class probabilities (log softmax) for img

```

```

probs = torch.nn.functional.softmax(model(imgs[0].unsqueeze(0)), dim=1)
# Plot the image and probabilities
view_prediction(imgs[0], labels[0], probs, dataset.classes)
break

```



### 3.13 Save the model

```

[44]: # Very important
torch.save(model, 'model')

```

### 3.14 Question 4 (10 points)

### 3.15 Compute accuracy on the test set

```

[45]: test_dataset = FashionMNIST('MNIST_data/', download = True, train = False,
↳ transform = transform)

```

```

[46]: test_dl = DataLoader(test_dataset, batch_size)
test_dl = DeviceDataLoader(test_dl, device)

```

```

[47]: def evaluate(model, test_dl):
      """
      Evaluates your model on the test data.

      Args:
          model: ImageClassifierNet object
          test_dl: test dataloader

```

```

Returns:
    Test accuracy.
    """
model.eval()  # Set the model to evaluation mode
correct = 0
total = 0

with torch.no_grad():  # Disable gradient computation
    for images, labels in test_dl:
        # Forward pass
        outputs = model(images)

        # Get predictions
        _, predicted = outputs.max(1)

        # Count correct predictions
        total += labels.size(0)
        correct += predicted.eq(labels).sum().item()

    # Calculate and return accuracy
    accuracy = correct / total
    return accuracy

```

```
[48]: print("Test Accuracy = {:.4f}".format(evaluate(model, test_dl)))
```

Test Accuracy = 0.9118

### 3.16 Tips to increase the test accuracy

- **Data augmentation:** Diversifies your training set and leads to better generalization
  - Flipping
  - Rotation
  - Shifting
  - Cropping
  - Adding noise
  - Blurring
- **Regularization:** Reduces overfitting on the training set
  - Early stopping
  - Dropout
  - $l_2$  regularization
  - Batch normalization
- **Hyperparameter tuning:**
  - Weight initialization
  - Learning rate
  - Activation functions
  - Optimizers